

Department of Computer Science



Submitted in part fulfilment for the degree of MEng.

Swarm Memory

Harry Burge

Version 1.0, 2020-November

Supervisor: Simon O'Keefe

Acknowledgements

Contents

Executive Summary	vi
1 Introduction	1
2 Literature Review	3
2.1 Cloud/Backup storage policys/schemes	3
2.2 Swarm robotics	5
3 Motiviation	9
4 Methodology/Design	10
4.1 The Problem	10
4.2 Simulation Infomation	11
4.3 Simple Scattered Memory Policy	11
5 Conclusion	14
A Some apendix	15
B Another apendix	16

List of Figures

2.1	Example of a hetrogenus ant colony. https://www.pinterest.co.uk/pin/77736358565153284	
4.1	Example of data density based of coordiante point	10
4.2	Example of simulation looks when running	12

List of Tables

Executive Summary

1 Introduction

Swarm intelligence/mechanics is an increasingly more important area of research for society, as the world moves towards a distributed technology future. Swarm intelligence can be viewed as distributed problem solving[2, 5], this is ever becoming more relevant as computer systems start to level out in terms of individual performance [7] and parallelism is embraced, satisfying the demand of the age of big data [9]. Swarm mechanics/robotics are on the rise in the industry, as society's pace increases and manual labor is automated out, whether its drone delivery to inpatient customers or mapping areas in dangerous environments [1].

An area of swarm intelligence research is distributed and local memory of swarm-like agents. This has gone down a route more to do with the optimization of distributed problem-solving algorithms rather than practical applications of storage of abstract ideas as a collective. There is a relative lack of research into collective memory on swarm like agents, within the scope of a practical setting. This would be invaluable to applications like mapping of a dangerous area [2], being able to handle the loss of agents and the collection of data on agents with limited memory. An explanation for this to be a less developed area of study is the existence of subjects like cloud-based and raid based storage systems.

Storing of data on an ever-changing network of storage devices is a hard task to complete, handling the loss of connection between different servers, reliability to access of the data and handling loss of services, whether it be non-correlated or correlated failures [9]. This is very applicable to swarm memory, with the handling of data across a swarm. However, must be adapted due to current algorithms such as RAID not being designed for highly dynamic systems such as swarm. There are promising papers in this field of cloud storage that suggest approaches that can be adapted with few modifications to apply to a swarm-based system [10], and optimized through RAID.

The objectives of this report are to merge three different areas of study into one by using knowledge of each, to create a suitable storage policy for swarm like agents to store collective memories of abstract ideas, then to perform analysis on a variety of simulations to explore the capability of the said storage policy.

1 Introduction

// Talk about the sections of the report once completed.

2 Literature Review

This chapter will review two areas of relevancy to the project proposed, these are Cloud/Backup storage policies/schemes and Swarm robotics. Ideas and concepts from both areas will have to be relied upon for the completion/design of the storage policy.

2.1 Cloud/Backup storage policys/schemes

Like most things in computer science, this area of study used to be simple, with small data-warehouses and backups on to a medium like magnetic tape, following a grandfather, father, son backup policy. As the years have progressed, technology has become greatly complex requiring larger files to be stored and accessed frequently. Leading to the need for complex backup systems to provide availability and longevity of data stored, across a network or even locally. A component to the lead of complexity of these algorithms other than providing a service better than competitors is the Legal Services Act. 2007 [11]. This enforces that a company's cloud storage solution has to be reliable and fast in data collection for users.

Most algorithms used in production are called random replication policies [9], this is where data is partitioned and randomly distributed among other storage devices usually on different racks of the datacenter. This is an efficient design policy for handling non-correlated errors, however, lacks the robustness against correlated errors. These algorithms are substantial for long-term data storage with average popularity of collection/use of that data.

A non-correlated error is when devices go down randomly for example in a swarm an agent's battery might explode, destroying that singular agent, therefore it can be modeled by a random chance of happening on each member. A correlated error happens when multiple agents go down due to a reason unbenounced to us.

Following on from our swarm example, let us say a tsunami hits a certain section of the swarm destroying those agents therefore meaning that those failures were connected by some sort of event. In data-warehouses, this is

2 Literature Review

usually a failure of power on a server or a bug, not as permanent damage as a tsunami.

Some of the problems arisen from random replication policy has spawned new replication policies, which can handle correlated and non-correlated errors, more effectively whilst also taking into account the demand of such items stored [9, 10]. Tackling these problems, two approaches were undertaken. The first approach is to come from a higher level of control where you have a manager which can choose items to replicate and where based on demand, knowledge of other replications and outside factors [9, 12]. This does not only apply to data but also schema changes of said servers or databases [13]. Working versions of these over a cloud service are tied in/together with a "Distributed key-value store" [14] where you have these key-value pairs on multiple devices on a network where duplication only leads to more fault tolerance of the data stored.

Another way to handle this which doesn't rely on a more privileged controller, and creates a distributed system is by having something like "SKUTE" as proposed in [10]. Each individual key-value replication has its own manager and can choose what it as a singular entity can do on that distributed system. The policy as described in "SKUTE" are as follows; Migration, Suicide, Replication, and Nothing. Migration is the moving of its data to lower costing and more redundant servers. Suicide is the removal of itself usually based on the number of duplicates and uses something like Paxos [15] to decide if it should suicide. Replication decides that it is being used enough to warrant the need to be duplicated.

An algorithm like "SKUTE" e.g. scattered key-value store [10, 16] will be best suited for swarm like agents due to the distributed nature of a swarm. This is less relevant in heterogeneous swarms with something like hivemind or hierarchal control however still relevant. For homogeneous swarms this is ideal due to not wanting to have a static/temporary leader because of the issues like communication bottleneck, power loss near the leader due to flow of information to the leader, loss of a leader in a hybrid static swarm [1, 5]. Also, one of the fundamental philosophies of swarm robotics is the inherently parallel nature that they bring to the table. When creating leader-based algorithms they do not fit within the spirit of the design of homogeneous swarms, unlike a heterogeneous swarm might.

An area of study which is stagnated is the storage of data on local disks. How to keep either backup for disk failure and/or improve write performance onto disks, rather than just duplicating data like as described above. An example schema for this is RAID, which has different levels based on the type of attributes that you may need for your array of storage devices [?]. In terms of cloud-based storage RAID arrays are used commonly internally rather than externally to a different NAS or storage server. This is because

we will have the guarantees of RAID for internal disk failures if they occur, for example, a server goes down we have the above replication schemes to be able to handle that loss of storage. Leading to RAID arrays across multiple nodes to be sort of redundant. However, by using a parity [8] based higher-level schema you can get space savings on the duplication of data. This is harder to implement though and most likely not needed due to servers that are lost due to power outages usually coming back online pretty soon. An example of this is might be a power issue causing a server to restart.

// TODO: Say what a correlated and non correlated failure is

2.2 Swarm robotics

Within the research of swarms, there is a split between practical solutions of agents or fictitious agents for an algorithm. This split can be seen as swarm robotics and swarm intelligence. Swarm intelligence is where we use distributed swarm-like behavior to solve a problem, for example traveling salesman problem [5], this means that we use an agent like code to compute the solution to a task. These tasks have usually been solved using a different algorithm beforehand, like TSP using a genetic algorithm to solve, and the swarm algorithms like AS-TSP [5] are alternatives to that algorithm. Said algorithms provide benefits and drawbacks compared to their counterparts. An area in which these algorithms could excel, and researched, is the networking space [5] due to the natural parallelism that can be exploited from the design of a distributed solver. The main concept of swarm intelligence is creating a solver to a problem using a distributed algorithm that doesn't rely on global knowledge and can be adaptable on the fly.

This is not the route of research that will be needed for this project, rather we will be looking at swarm robotics research. Swarm robotics has the same concept as swarm intelligence, it focuses on tasks that are designed to have agents complete, mainly for the practical space like moving objects or mapping an area [2, 4], whether simulated or not. Swarm robotics focuses on the behavior of the swarm more than the solution that it gives us. These swarms come in three types: heterogeneous, homogenous, and a subcategory of homogeneous, hybrid swarms [1]. These three methodologies are mapped onto both the decision making and the agent's body/abilities.

A heterogeneous swarm is where there are differences between the agents, as in Figure 2.1, whether physical or mental [1, 5]. These most commonly

2 Literature Review

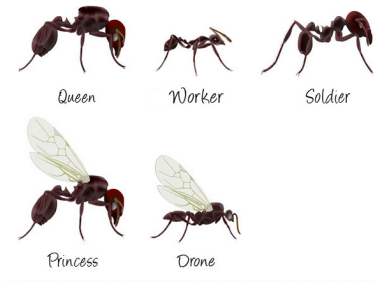


Figure 2.1: Example of a heterogeneous ant colony. <https://www.pinterest.co.uk/pin/777363585651532845/>

occur in nature and are not usually studied [5] due to the differences in agents being a rarely needed property in research-based problems. In real-world solutions, heterogeneous swarms can be of great use, allowing other agents to pick up the slack of the swarm, or complete tasks that other swarm members cannot complete. For example, as described in [1] with a mother ship being a navy boat and a swarm of quadcopters, this example shows how the boat picks up for the slack of the swarm by being able to transport them longer distances than the swarm could normally cope with. Less research is done into this area is due to some key drawbacks of having a heterogeneous swarm. Usually, you will have a hive-mind like a system if you have a heterogeneous swarm where you have leaders giving commands to subordinates or even one leader commanding the entire swarm. This is less desirable if there is a loss of those leaders you lose the ability to control the swarm, in our example, this doesn't matter so much due to if you have a loss of the mothership something has gone significantly wrong already. . Due to the differences in the swarm agents, it allows for greater efficiency of the swarm, having robots that can mine and that can farm exclusively. This leads to the vulnerability of major loss of one type of agent can lead to the loss of the colony.

With swarms like these to get the efficiency from a heterogeneous swarm without leading to vulnerabilities, you need to have some jobs be interchangeable, e.g. agents are adaptable like in an ant colony [22]. Also, the jobs that can't be done by all agents need not be vital to the survivability of the swarm. Ants usually fit into this type of swarm where you have a queen, worker, and major ants, some ants like Leaf-Cutter Ants also have a subcategory of workers like a fungus farmer. If there is a significant loss of workers, majors start doing tasks that normally workers would do [5]. With the farmer subcategory, if a significant loss happens other workers/majors can take over that job and learn how to do it. This leads us more towards homogeneous agents within a heterogeneous swarm. This is one of the main reasons for not wanting to use heterogeneous swarms

2 Literature Review

because human-made machines are less adaptable in this job switching method, without redundant hardware or software.

A homogeneous swarm is where each agent is the same, found less often in nature, and is more towards man-made agents. This is due to nature taking to a more efficient approach, having variety in a semi-homogenous swarm allows for greater efficiency, and being of a natural organism has more adaptability than a man-made robot agent [1, 5]. In homogeneous swarms we get significant redundancy, if an agent goes down we have swarms worth of replacements for that agent. With this redundancy, we gain possible losses of efficiency, either an agent is too simple therefore losing specialism or each agent has parts for specialism but may never need said part. An example of this homogenous flaw is we have a humanoid agent that has arms, and we want the swarm of agents to mine and farm. If the humanoid has arms only then it will take them longer to do these tasks, compared to having specific equipment for the job. Flipping this the other way around, if each agent has equipment for mining and farming, some of these agents won't need both items and means the swarm would use more resources to create. In a homogeneous style of swarm, we have homogeneous control, also known as distributed intelligence, where all agents decide what they want to do based on what other agents are doing around them, an easy to understand example of this is [23]. This decision making involves internal parameters and can be equated to an emergent/structured swarm. Homogeneous control is different from heterogeneous control due to following a distributed problem solving/communication design compared to leader based design like hivemind or structured/hierarchical controlled swarms.

Within swarm robotics everything gets a bit messy, usually, there is no clear-cut name or design that can be assigned to swarm models and behaviors. This is down to wanting the agents to act in a specific emergent pattern rather than being stuck to a strict certain arbitrary rule, this is where hybrid approaches come into play. A hybrid approach is a mixture of both heterogeneous and homogeneous natures in both communications and agent design [1, 5]. In terms of communication when taking a hybrid approach to a swarm, you will have a swarm leader/leaders designated by the swarm, and handled with a consensus algorithm, e.g. Paxos [15]. In a hybrid model for agent design, if we want to gain the efficiency of the heterogeneous model and the adaptability/reliability of a homogeneous swarm we can use something like part-time tools. This allows our homogeneous bots to act in heterogeneous fashions.

Humans, themselves are a great example of a hybrid based swarm both in design and communication. Though humans have variations in characteristics they can be seen as pretty homogeneous in terms of the tasks that they can perform, obviously removing edge case actions that humans

2 Literature Review

do. Tools and knowledge can be spread between humans to make the swarm more efficient and an agent can specialize in a certain area however if some agents are lost other agents/humans can replace them by using the same tools and learning from the remaining agents of that task. Also, the natural power-based structure of humans fits a hybrid model in terms of electorship of some kind, and not of genetics (Except with royalty, however, this is more of a label rather than a genetic difference). The leaders aren't needed for every single action so fit into a usually hierarchal power structure, compared to something of a hivemind model.

3 Motivation

As described above in Chapter 2 it is quite clear that swarm robotics is becoming and already is a solution to many problems currently faced and going on into the future. With individual speeds of computers stagnating [7], the world becoming more data oriented and the drive for humans to explore, distributed technologies need to rise to the challenge. The next big step in computer science history being quantum computers, not solving the sequential issues of our current problem solving algorithms. This is why swarm robotics/intelligence needs to be brought to the forefront of research.

In our day and age there doesn't seem to be much use of these innovative works, possibly due to the complexity or to the newness of said subject. The main applications nowadays are surveillance [21, 18] and delivery, however looking into the future, which we need to do otherwise we delay technological breakthroughs further, we can see that there are so many other uses for this type of distributed thinking. Whether it be physically with robots, like in space exploration [20], nano-robot medicine or military based applications. Or even conceptually like algorithms that can rely more on parallel computation compared to their global sequential counterparts. A good example of this distributed thinking changing and revolutionising our subject is block chain [19].

It is for this reason that I have decided to contribute my part to this extensive and breakthrough field.

4 Methodology/Design

4.1 The Problem

The problem is to create an efficient storage system for directional based data. This can be broken down into two sections, how data is passed throughout the swarm as to not have too many duplicates of the data but not have too little to not allow for redundancy. The second is to have that data spread out based on a coord point, this will simulate the use of said data in that area.




Figure 4.1: Example of data density based on coordinate point

As shown in Figure 4.1 we can see how we want the data to spread out across the swarm based off the distance of the agent to the point. An example of why this is a requirement is for example in that arbitrary point might be the farm and possible hostilities around, therefore you want more agents knowing that arbitrary point information around that area but still have that information spread across the network due to possible losses of all the agents near said point.

The goal is to create an efficient storage and replication policy which can handle the fluctuations of the swarm, and provide reliability in the face of both correlated failures and non-correlated failures of the swarm. This policy should take into account position of the agent relative to the point that the information is needed near.

4.2 Simulation Information

The base simulation that the policies will be running on is a specified. There will be an area of which agents can move around, each point within that area can have a piece of data learned in that area. When a piece of data is learned in that area, the data stores where it was created and the policies will take that starting place into account. This is how we have defined our directional data aspect to the simulation, in real life this could be something like a learnt behaviour that is needed for that specific environment, however, for test purposes it is just a string.

Each agent has two bounded partitions of memory, these are called private and public memory. Private memory is for knowledge that is learnt by that agent, and public memory is information that that agent has learned from other agents. Both public and private memory cannot have duplicate data inside including across partitions. This split was made arbitrarily and isn't needed for a practical solution however was created due to agents being more likely to pick up learned data by not having to delete information in memory to make space for said learned information. In general public memory should be larger in size than private memory due to there being magnitudes more public information compared to private information for an agent.

Agents are homogenous, and have the same connection radius of which they can talk to other agents within. Agents are assumed to be able to receive and send packets whilst also doing actions like sending and receiving itself, in real world solutions this would likely be that each agent would need two receivers and senders or a control algorithm to handle this type of behaviour.

Each agent moves randomly in circles varying in shapes and sizes to gain a swarm like movement of connections being broken and reconnected. Agents are running in different threads so run at different speeds and act more like a real world swarm. Data is introduced at the start based on an amount wanted, and during the simulation data can be learned based on an amount wanted.

4.3 Simple Scattered Memory Policy

Taking on from the design of algorithm presented in [10], I applied a simple distributed storage policy onto a swarm that takes into account positional data. This policy has two actions that it can perform, suicide and replication,

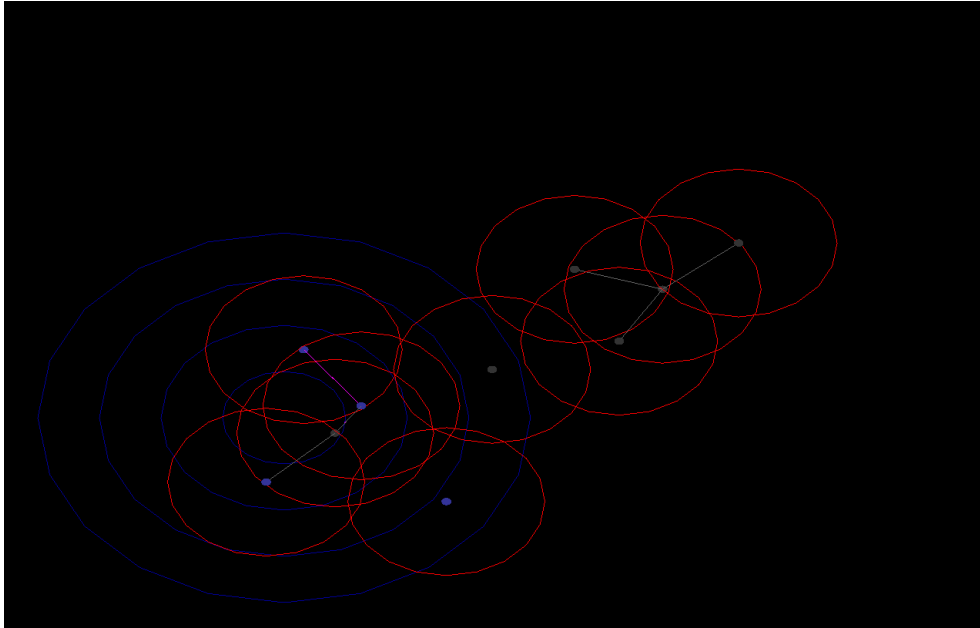


Figure 4.2: Example of simulation looks when running

suicide is when a piece of data decides using a heuristic that it is not worth being stored in an agents memory so deletes itself, replication is where based on a heuristic the data believes it is worth spreading the information so will replicate to one member of the swarm.

Algorithm 1 Agent's control loop

```

1: procedure STEP
2:   move()
3:
4:   if item in private mem not in other agents public mem then
5:     Replicate item
6:     return true
7:
8:    $item \leftarrow$  random public mem item
9:    $allowedupes \leftarrow heuristic(item)$ 
10:   $currentdupes \leftarrow$  number of agents with item in public mem
11:
12:  if allowed dupes > current dupes then
13:    Replicate item once
14:  else if allowed dupes < current dupes then
15:    Suicide item using paxos
16:
17:  return true
18:
19: procedure HEURISTIC(ITEM)
20:   $dist \leftarrow$  agent distance to items target point
21:   $allowedupes \leftarrow$  items max dupes  $- (dist/items\ step)$ 
22:
23:  if allowed dupes  $\leq 0$  then
24:    return 0
25:  else
26:    return allowed dupes

```

5 Conclusion

A Some apendix

B Another apendix

Bibliography

- [1] J. C. Barca and Y. A. Sekercioglu, "Swarm robotics reviewed," *Robotica*, vol. 31, no. 3, pp. 345–359, 2013.
- [2] V. Kumar and F. Sahin, "Cognitive maps in swarm robots for the mine detection application," *SMC'03 Conference Proceedings. 2003 IEEE International Conference on Systems, Man and Cybernetics. Conference Theme - System Security and Assurance (Cat. No.03CH37483)*, Washington, DC, 2003, pp. 3364-3369 vol.4, doi: 10.1109/ICSMC.2003.1244409.
- [3] H. Wang, D. Wang and S. Yang, "Triggered Memory-Based Swarm Optimization in Dynamic Environments," in *Applications of Evolutionary Computing*, M. Giacobini, Ed. Berlin, Germany: Springer-Verlag Berlin and Heidelberg GmbH & Co. K, 2007, pp. 637–646.
- [4] D. A. Lima and G. M. B. Oliveira, "A probabilistic cellular automata ant memory model for a swarm of foraging robots," *2016 14th International Conference on Control, Automation, Robotics and Vision (ICARCV)*, Phuket, 2016, pp. 1-6, doi: 10.1109/ICARCV.2016.7838615.
- [5] E. Bonabeau, M. Dorigo, and G. Theraulaz, *Swarm Intelligence: From Natural to Artificial Systems*. Cary, NC, USA: Oxford University Press, 1999.
- [6] L. Xiang, Y. Xu, J. Lui, Q. Chang, Y. Pan, and R. Li, 'A Hybrid Approach to Failed Disk Recovery Using RAID-6 Codes: Algorithms and Performance Evaluation', *Association for Computing Machinery*, vol. 7, p. 11, 2011
- [7] C. Mims, 'Why CPUs Aren't Getting Any Faster', *MIT Technology Review*, 2010. [Online]. Available: <https://www.technologyreview.com/2010/10/12/199966/why-cpus-arent-getting-any-faster/>. [Accessed: 01-Dec-2020].
- [8] U. Troppens, W. Müller-Friedt, R. Wolafka, R. Erkens, and N. Haustein, 'Appendix A: Proof of Calculation of the Parity Block of RAID 4 and 5', in *Storage Networks Explained: Basics and Applic-*

Bibliography

- ation of Fibre Channel SAN, NAS, iSCSI, InfiniBand and FCoE, U. Troppens, Ed. Chichester: Wiley United Kingdom, 2009, pp. 535–536.
- [9] J. Liu and H. Shen, "A Low-Cost Multi-failure Resilient Replication Scheme for High Data Availability in Cloud Storage," 2016 IEEE 23rd International Conference on High Performance Computing (HiPC), Hyderabad, 2016, pp. 242-251, doi: 10.1109/HiPC.2016.036.
- [10] N. Bonvin, T. G. Papaioannou, and K. Aberer, A Self-Organized, Fault-Tolerant and Scalable Replication Scheme for Cloud Storage. New York, NY, USA: Association of Computing Machinery, 2010.
- [11] Legal Services Act. 2007.
- [12] A. Prahlad, M. S. Muller, R. Kottomtharayil, S. Kavuri, P. Gokhale, and M. Vijayan, 'Cloud gateway system for managing data storage to cloud storage sites', 20100333116A1, 2010.
- [13] B. Czejdo, K. Messa, T. Morzy, M. Morzy, and J. Czejdo, 'Data Warehouses with Dynamically Changing Schemas and Data Sources', in Proceedings of the 3rd International Economic Congress, Opportunities of Change, Sopot, Poland, 2003, p. 10.
- [14] 'Key-Value Scores Explained', HazelCast. [Online]. Available: <https://hazelcast.com/glossary/key-value-store/>. [Accessed: 02-Dec-2020].
- [15] L. Lamport, 'The Part-Time Parliament', in Concurrency: The Works of Leslie Lamport, New York, NY, USA: Association of Computing Machinery, 2019, pp. 277–317.
- [16] D. Agrawal and A. E. Abbadi. The tree quorum protocol: An efficient approach for managing replicated data. In VLDB'90: Proc. of the 16th International Conference on Very Large Data Bases, pages 243–254, Brisbane, Queensland, Australia, 1990.
- [17] S. Lynn, 'RAID Levels Explained', PC Mag, 2014. [Online]. Available: <https://uk.pcmag.com/storage/7917/raid-levels-explained>. [Accessed: 06-Dec-2020].
- [18] J. Hu et al., Eds., HiveMind: A Scalable and Serverless Coordination Control Platform for UAV Swarms. ArXiv, 2020.
- [19] D. Calvaresi, A. Dubovitskaya, J. P. Calbimonte, K. Taveter, and M. Schumacher, Multi-Agent Systems and Blockchain: Results from a Systematic Literature Review. Cham, Switzerland: Springer Interna-

Bibliography

tional Publishing, 2018.

- [20] L. A. Nguyen, T. L. Harman and C. Fairchild, "Swarmathon: A Swarm Robotics Experiment For Future Space Exploration," 2019 IEEE International Symposium on Measurement and Control in Robotics (IS-MCR), Houston, TX, USA, 2019, pp. B1-3-1-B1-3-4, doi: 10.1109/IS-MCR47492.2019.8955661.
- [21] M. Y. Arafat and S. Moh, "Localization and Clustering Based on Swarm Intelligence in UAV Networks for Emergency Communications," in IEEE Internet of Things Journal, vol. 6, no. 5, pp. 8958-8976, Oct. 2019, doi: 10.1109/JIOT.2019.2925567.
- [22] D. Jackson and F. Ratnieks, 'Communication in ants,'Current Biology,vol. 16, pp. 570–574, 2006.
- [23] C. W. Reynolds, Flocks, Herds, and Schools: A Distributed Behavioral Model. ACM, 1987.